

ISYE-6420 Midterm Exam

Marshall Bockrath-Vandegrift

March 20, 2016

Contents

1	Master of light	2
3	Amanita muscaria	7
4	Squids	11

1 Master of light

1 (a) One-way Bayesian ANOVA

We will model each measurement as $X \sim \mathcal{N}(\mu_0 + \alpha_{\text{type}}, \sigma^2)$, with improper flat priors across the support of all latent parameters and ensuring identifiability by constraining the α parameters to sum to zero. We implement this model in Stan,

```
data {
  int<lower=1> N;
  int<lower=1> NTYPE;
  int<lower=1, upper=NTYPE> type[N];
  vector[N] meas;
}
parameters {
  vector[NTYPE-1] alpha_raw;
  real mu0;
  real<lower=0> sigma;
}
transformed parameters {
  vector[NTYPE] alpha;
  for (i in 1:(NTYPE-1)) alpha[i] <- alpha_raw[i];
  alpha[NTYPE] <- -sum(alpha_raw);
}
model {
  meas ~ normal(mu0 + alpha[type], sigma);
}
generated quantities {
  real diff[NTYPE, NTYPE];
  for (i in 1:NTYPE)
    for (j in 1:NTYPE)
      diff[i, j] <- alpha[i] - alpha[j];
}
```

We generate samples of the α parameters from this model, yielding the following summary:

Inference for Stan model: pla-anova-1way.

4 chains, each with iter=49000; warmup=24500; thin=1;

post-warmup draws per chain=24500, total post-warmup draws=98000.

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
alpha[1]	15.26	0.02	4.11	7.16	12.51	15.26	18.03	23.28	63321	1
alpha[2]	-0.91	0.01	3.77	-8.36	-3.44	-0.92	1.61	6.49	65039	1

alpha[3]	-1.45	0.01	2.89	-7.11	-3.38	-1.47	0.48	4.23	63511	1
alpha[4]	-4.29	0.01	3.68	-11.51	-6.76	-4.31	-1.82	2.98	65455	1
alpha[5]	-8.61	0.01	3.96	-16.41	-11.23	-8.60	-5.97	-0.84	98000	1

Samples were drawn using NUTS(diag_e) at Fri Mar 18 08:56:52 2016.
For each parameter, n_eff is a crude measure of effective sample size,
and Rhat is the potential scale reduction factor on split chains (at
convergence, Rhat=1).

We also plot the 95% equi-tailed credible sets for the α parameters (figure 1).
As seen in both the numeric summary and the plot, it appears that the treatment
effects alpha[1] (“Glass/8”) and alpha[5] (“Steel/12”) differ statistically from the
overall mean measurements.

1 (b) One-way treatment comparisons

Samples of the differences between treatment effects (diff in the Stan model) yield
the following summary:

Inference for Stan model: pla-anova-1way.
4 chains, each with iter=49000; warmup=24500; thin=1;
post-warmup draws per chain=24500, total post-warmup draws=98000.

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
diff[1,2]	16.17	0.03	6.35	3.65	11.94	16.15	20.43	28.65	60502	1
diff[1,3]	16.71	0.02	5.54	5.82	13.00	16.72	20.44	27.55	59758	1
diff[1,4]	19.55	0.03	6.27	7.19	15.35	19.55	23.76	31.85	61063	1
diff[1,5]	23.87	0.02	6.54	10.97	19.52	23.85	28.26	36.73	98000	1
diff[2,3]	0.54	0.02	5.13	-9.56	-2.88	0.56	3.95	10.62	63847	1
diff[2,4]	3.38	0.02	5.89	-8.26	-0.57	3.40	7.35	14.89	60300	1
diff[2,5]	7.69	0.02	6.20	-4.51	3.55	7.70	11.84	19.89	98000	1
diff[3,4]	2.84	0.02	5.00	-7.04	-0.51	2.85	6.18	12.67	62641	1
diff[3,5]	7.15	0.02	5.33	-3.33	3.61	7.15	10.69	17.72	98000	1
diff[4,5]	4.32	0.02	6.12	-7.70	0.24	4.31	8.40	16.35	98000	1

Samples were drawn using NUTS(diag_e) at Fri Mar 18 08:56:52 2016.
For each parameter, n_eff is a crude measure of effective sample size,
and Rhat is the potential scale reduction factor on split chains (at
convergence, Rhat=1).

We also plot the 95% equi-tailed credible sets for these differences (figure 2).
Both the plot and numeric summary suggest that only the “Glass/8” mirror type
showed a significant difference, producing statistically larger values for the precision
measurements versus all the other types.




Figure 1: Mirror type effect credible sets




Figure 2: Mirror type effect difference credible sets

1 (c) Two-way Bayesian ANOVA

We will define a similar model for the two-way problem, this time considering each measurement to be $X \sim \mathcal{N}(\mu_0 + \alpha_{\text{mat}} + \beta_{\text{sid}} + \gamma_{\text{mat},\text{sid}}, \sigma^2)$. In Stan,

```
data {
  int<lower=1> N;
  int<lower=1> NMAT;
  int<lower=1> NSID;
  int<lower=1, upper=NMAT> mat[N];
  int<lower=1, upper=NSID> sid[N];
  vector[N] meas;
}
parameters {
  vector[NMAT-1] alpha_raw;
  vector[NSID-1] beta_raw;
  matrix[NMAT-1, NSID-1] gamma_raw;
  real mu0;
  real<lower=0> sigma;
}
transformed parameters {
  vector[NMAT] alpha;
  vector[NSID] beta;
  matrix[NMAT, NSID] gamma;
  for (i in 1:(NMAT-1)) alpha[i] <- alpha_raw[i];
  alpha[NMAT] <- -sum(alpha_raw);
  for (i in 1:(NSID-1)) beta[i] <- beta_raw[i];
  beta[NSID] <- -sum(beta_raw);
  for (i in 1:(NMAT-1))
    for (j in 1:(NSID-1))
      gamma[i, j] <- gamma_raw[i, j];
  for (i in 1:(NMAT-1))
    gamma[i, NSID] <- -sum(gamma_raw[i, :]);
  for (j in 1:(NSID-1))
    gamma[NMAT, j] <- -sum(gamma_raw[:, j]);
  gamma[NMAT, NSID] <- sum(gamma_raw);
}
model {
  vector[N] mu;
  for (i in 1:N)
    mu[i] <- mu0 + alpha[mat[i]] + beta[sid[i]] + gamma[mat[i], sid[i]];
  meas ~ normal(mu, sigma);
}
```




Figure 3: Material-sides interaction effect credible sets

We generate samples of the γ parameters from this model, yielding the following summary:

Inference for Stan model: plc-anova-2way.

4 chains, each with iter=49000; warmup=24500; thin=1;

post-warmup draws per chain=24500, total post-warmup draws=98000.

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
gamma[1,1]	2.96	0.01	1.87	-0.74	1.73	2.97	4.21	6.63	72041	1
gamma[1,2]	-2.96	0.01	1.87	-6.63	-4.21	-2.97	-1.73	0.74	72041	1
gamma[2,1]	-2.96	0.01	1.87	-6.63	-4.21	-2.97	-1.73	0.74	72041	1
gamma[2,2]	2.96	0.01	1.87	-0.74	1.73	2.97	4.21	6.63	72041	1

Samples were drawn using NUTS(diag_e) at Fri Mar 18 08:57:05 2016.

For each parameter, n_eff is a crude measure of effective sample size, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat=1).

We also plot the 95% equi-tailed credible sets for the γ parameters (figure 3). As seen in both the numeric summary and the plot, the interaction effect does not meet our 95% confidence requirement for significance, but comes fairly close. More investigation may be necessary.

3 Amanita muscaria

3 (a) Gibbs sampler with gamma prior on τ

The distributions specified for the problem model lead to the standard conjugate relationships and marginal distributions.

```
sampler_gibbs_g <- function(x=spores,
                             mu0=12, sigma0=2, alpha=2, beta=4,
                             mu=NULL, tau=NULL,
                             warmup=1000, keep=100000) {
  tau0 <- sigma0 ^ -2
  n <- length(x)
  sum_x <- sum(x)
  rmu <- function(tau) {
    mean <- (tau0 * mu0 + tau * sum_x) / (tau0 + n * tau)
    sd <- (tau0 + n * tau) ^ -0.5
    rnorm(1, mean=mean, sd=sd)
  }
  rtau <- function(mu) {
    shape <- alpha + (n / 2)
    rate <- beta + sum((x - mu) ^ 2) / 2
    rgamma(1, shape=shape, rate=rate)
  }
  if (is.null(mu)) mu <- rnorm(1, mean=mu0, sd=sigma0)
  if (is.null(tau)) tau <- rgamma(1, shape=alpha, rate=beta)
  rtarget <- function() {
    mu <- rmu(tau)
    tau <- rtau(mu)
    c(mu=mu, tau=tau)
  }
  chain <- iter(rtarget)
  for (i in 1:warmup) nextElem(chain)
  do.call(rbind, as.list(ilimit(chain, keep)))
}
```

This sampler yields squared error loss point estimates of $\mu \approx 10.12$ and $\tau \approx 0.46$.

3 (b) Gibbs sampler with inverse gamma prior on σ^2

The distributions for this problem model are also standard conjugate, with the resulting standard marginal distributions.

```
sampler_gibbs_ig <- function(x=spores,
                             mu0=12, sigma0=2, alpha=4, beta=2,
                             mu=NULL, sigma2=NULL,
                             warmup=1000, keep=100000) {
  sigma02 <- sigma0 ^ 2
  n <- length(x)
  sum_x <- sum(x)
  rmu <- function(sigma2) {
    mean <- ((mu0 / sigma02 + sum_x / sigma2) /
              (1 / sigma02 + n / sigma2))
    sd <- (1 / sigma02 + n / sigma2) ^ -0.5
    rnorm(1, mean=mean, sd=sd)
  }
  rsigma2 <- function(mu) {
    shape <- alpha + (n / 2)
    rate <- beta + sum((x - mu) ^ 2) / 2
    1 / rgamma(1, shape=shape, rate=rate)
  }
  if (is.null(mu)) mu <- rnorm(1, mean=mu0, sd=sigma0)
  if (is.null(sigma2))
    sigma2 <- 1 / rgamma(1, shape=alpha, rate=beta)
  rtarget <- function() {
    mu <- rmu(sigma2)
    sigma2 <- rsigma2(mu)
    c(mu=mu, tau=(1 / sigma2))
  }
  chain <- iter(rtarget)
  for (i in 1:warmup) nextElem(chain)
  do.call(rbind, as.list(ilimit(chain, keep)))
}
```

The samples from this second model yield squared error loss point estimates of $\mu \approx 10.12$ and $\tau \approx 0.52$.

These results differ from those of the first model due to the change in hyperparameters. Repeating with the original hyperparameters but the second model's inverse gamma prior on σ^2 , we calculate our new estimates to be $\mu \approx 10.12$ and $\tau \approx 0.46$. These coincide with our estimates from the first model, reflecting the identical likelihoods resulting from the simple parameterization difference between the two models.

3 (c) Metropolis-Hastings sampler with gamma prior on τ

To produce a Metropolis-Hastings sampler we will need the density of our joint target distribution up to a normalizing constant. We can calculate the contribution of the normal density $p(\mathbf{x}|\mu, \tau)$ directly as $\prod_{i=1}^n p(x_i|\mu, \tau)$, but this is inefficient. We know from standard results in statistical theory that for the normal distribution \bar{X} is a sufficient statistic for location parameter μ , S^2 is sufficient for precision parameter τ , and by Basu's theorem \bar{X} and S^2 are independent. This allows us to calculate the density for the sample as the product of the densities of $\bar{X} \sim \mathcal{N}(\mu, [n\tau]^{-1})$ and $(n-1)S^2 \sim \tau^{-1}\chi^2(n-1)$. For the density of $(n-1)S^2$ we must include the contribution of the non-constant transform of the χ^2 random variable via the transform inverse absolute derivative $|(g^{-1})'(x)| = \tau$.

For our proposal distribution we will take uncorrelated Gaussian steps in both dimensions, making our sampler a Metropolis random walk sampler. This can generate proposed values outside of the support of τ , but we will assign such values a density of 0 and reject them as part of the normal proposal-evaluation process.

```
sampler_mh <- function(x=spores,
                        mu0=12, sigma0=2, alpha=2, beta=4,
                        mu=NULL, tau=NULL,
                        warmup=10000, keep=100000) {
  n <- length(x)
  x_bar <- mean(x)
  x_ssd <- (n - 1) * var(x)
  dtarget_log <- function(theta) {
    mu <- theta[1]; tau <- theta[2]
    if (tau <= 0) {
      -Inf
    } else {
      dnorm(x_bar, mean=mu, sd=((n * tau) ^ -0.5), log=TRUE) +
        dchisq(tau * x_ssd, df=(n - 1), log=TRUE) + log(tau) +
        dnorm(mu, mean=mu0, sd=sigma0, log=TRUE) +
        dgamma(tau, shape=alpha, rate=beta, log=TRUE)
    }
  }
  rproposal <- function(theta) {
    xi <- rnorm(length(theta), mean=theta, sd=c(1.5, 0.1))
    names(xi) <- names(theta)
    xi
  }
  if (is.null(mu)) mu <- rnorm(1, mean=mu0, sd=sigma0)
  if (is.null(tau)) tau <- rgamma(1, shape=alpha, rate=beta)
  theta <- c(mu=mu, tau=tau)
```

```

rtarget <- function() {
  xi <- rproposal(theta)
  rho <- exp(min(0, dtarget_log(xi) - dtarget_log(theta)))
  theta <- if (runif(1) <= rho) xi else theta
  theta
}
chain <- iter(rtarget)
for (i in 1:warmup) nextElem(chain)
do.call(rbind, as.list(ilimit(chain, keep)))
}

```

This sampler also yields the expected squared error loss point estimates of $\mu \approx 10.12$ and $\tau \approx 0.46$.

4 Squids

4 (a) Linear regression

We define a linear regression model in Stan, accounting for the assumed-MCAR x_1 and y values, supporting sampling-time prediction for new inputs, and calculating the Ibrahim-Laud L criteria and the predictive log-likelihood necessary for LOO estimation.

Parameter and predictive point estimates for this model show little sensitivity to priors, but their credible sets do, under entirely non-informative priors generating nonsensically wide intervals (e.g., including negative squid weights). Following advice in the Stan manual and Gelman et al. (2008),¹ we regularize the regression coefficients by centering and scaling each of the x variables to a common range then placing on the coefficients a common Cauchy prior scaled to the distribution of y .

This results in the following model, which while somewhat complicated we will use without further modification for the remaining parts of this problem.

```
functions {  
  matrix rescale(matrix x, row_vector x_bar, row_vector s) {  
    return ((x - rep_matrix(x_bar, rows(x)))  
            ./ rep_matrix(s, rows(x)) ./ 2);  
  }  
}  
  
data {  
  int<lower=1> K;  
  int<lower=1> N_cc;  
  int<lower=0> N_mis_x1;  
  int<lower=0> N_mis_y;  
  int<lower=0> N_pred;  
  row_vector[K] x_bar;  
  row_vector[K] s;  
  matrix[N_cc, K] x_cc;  
  matrix[N_mis_x1, K] x_mis_x1;  
  matrix[N_mis_y, K] x_mis_y;  
  matrix[N_pred, K] x_pred;  
  vector[N_cc] y_cc;  
  vector[N_mis_x1] y_mis_x1;  
}  
  
transformed data {  
  int N_obs;  
  int P;
```

¹Gelman, A., Jakulin, A., Pittau, M. G., and Su, Y.-S. (2008). A weakly informative default prior distribution for logistic and other regression models. *Annals of Applied Statistics*, 2(4):1360–1383.

```

matrix[N_cc, K] x_cc_std;
matrix[N_mis_x1, K] x_mis_x1_std;
matrix[N_mis_y, K] x_mis_y_std;
matrix[N_pred, K] x_pred_std;
vector[N_cc + N_mis_x1] y_obs;
real scale_beta;
real sst;
N_obs <- N_cc + N_mis_x1;
P <- K + 1;
x_cc_std <- rescale(x_cc, x_bar, s);
x_mis_x1_std <- rescale(x_mis_x1, x_bar, s);
x_mis_y_std <- rescale(x_mis_y, x_bar, s);
x_pred_std <- rescale(x_pred, x_bar, s);
y_obs <- append_row(y_cc, y_mis_x1);
scale_beta <- 2.5 * sd(y_obs);
sst <- dot_self(y_obs - mean(y_obs));
}
parameters {
  real<lower=0> sigma;
  real beta0;
  vector[K] beta;
  vector[N_mis_y] y_mis_y;
  vector[N_mis_x1] x1_imputed;
}
transformed parameters {
  matrix[N_obs, K] x_obs_std;
  x_obs_std <- append_row(x_cc_std, x_mis_x1_std);
  if (N_mis_x1 > 0)
    x_obs_std[(N_cc+1):, 1] <- (x1_imputed - x_bar[1]) / s[1] / 2;
}
model {
  beta ~ cauchy(0, scale_beta);
  if (N_mis_x1 > 0) x1_imputed ~ normal(3.0/2.0, 0.5);
  if (N_mis_y > 0)
    y_mis_y ~ normal(beta0 + x_mis_y_std * beta, sigma);
  y_obs ~ normal(beta0 + x_obs_std * beta, sigma);
}
generated quantities {
  vector[N_obs] log_lik;
  vector[N_pred] y_pred;
  real BR2;
  real L2;

```

```

{
  vector[N_obs] Z;
  for (i in 1:N_obs) {
    real mu;
    mu <- beta0 + x_obs_std[i] * beta;
    log_lik[i] <- normal_log(y_obs[i], mu, sigma);
    Z[i] <- normal_rng(mu, sigma);
  }
  L2 <- dot_self(Z - y_obs);
  BR2 <- 1 - (N_obs - P) * sigma^2 / sst;
  for (i in 1:N_pred)
    y_pred[i] <- normal_rng(beta0 + x_pred_std[i] * beta, sigma);
}
}

```

Using the resulting samples, we estimate this model to have a Bayesian R^2 of 0.95. We estimate the missing $x_1 \approx 1.38$ and the missing $y \approx 9.39$.

4 (b) Prediction

Our model predicts y for the specified x values to be 3.29 with squared error loss, and a 95% equi-tailed credible set of (0.1, 6.49).

4 (c) Model selection

We re-fit the model five times, each time removing one of the x predictor variables and recording the resulting L^2 and log-likelihood values. We then use these values to compute for each reduced model the Ibrahim-Laud $L = \sqrt{L^2}$ criteria and the deviance-scale PSIS-LOO² estimated leave-one-out cross validation expected log-likelihood.

	full	x_1	x_2	x_3	x_4	x_5
L	4.85	4.74	4.99	4.72	5.19	6.06
looic	57.35	55.00	55.20	53.94	59.32	66.02

In this case both metrics agree, taking lowest values for the model eliminating variable x_3 .

For these posterior predictive checks, we have elected to include the observation missing x_1 under multiple imputation, but exclude the observation missing y . It does not seem to make sense to check the performance of a model versus an observation predicted via the model. That said, we generated the predictive performance metrics under each combination of including/excluding the missing-data observations and found the same model (that excluding x_3) most effective under all combinations.

²Vehtari, A., Gelman, A., and Gabry, J. (2016). Practical Bayesian model evaluation using leave-one-out cross-validation and WAIC. arXiv preprint: <http://arxiv.org/abs/1507.04544>